

Quantum-enhanced machine learning: Implementing a quantum k -nearest neighbour algorithm

Bachelor thesis defense

19. January 2017

Mark Fingerhuth

Maastricht Science Programme, Maastricht University, The Netherlands
Thesis work at the Centre for Quantum Technology, UKZN, South Africa




Table of contents

1. Introduction
2. Machine Learning
3. Quantum Computing
4. Quantum-enhanced Machine Learning
5. Results: Qubit-based kNN algorithm
6. Results: Amplitude-based kNN algorithm
7. Conclusion

Introduction

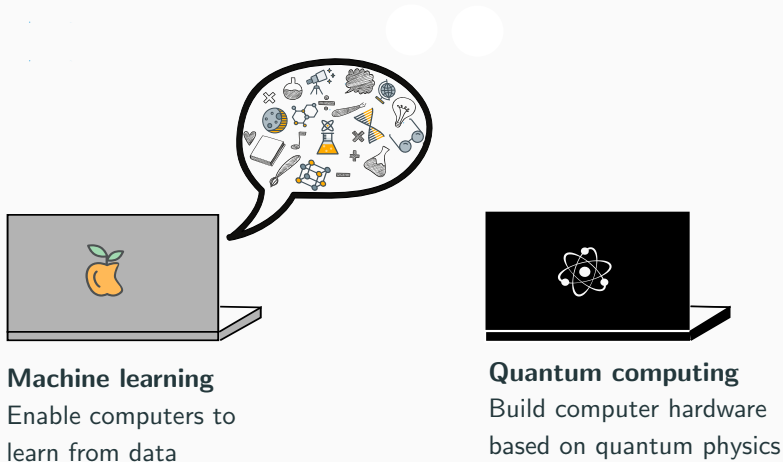
Enhancing machine learning with quantum mechanics



Machine learning

Enable computers to
learn from data


Enhancing machine learning with quantum mechanics



Machine Learning


Supervised pattern recognition

Introduce the concept of training and input dataset labels as well as supervision




k -nearest neighbour algorithm

$k = 3$



$k = 7$



k -nearest neighbour algorithm

Explain the kNN!

Quantum Computing

Quantum bits (Qubits)

0 and 1 introduce ket vectors & superposition show vector representation of a qubit

Visualizing a qubit

Bloch sphere intro maybe not needed if I concentrate on qubit-based kNN only

show matrix representation show matrix vector multiplication?? Show two simple examples (Hadamard and X gate?)

Quantum-enhanced Machine Learning

Encoding classical data into qubits

Encoding classical data into amplitudes

Amplitude interference

Results: Qubit-based kNN algorithm

Results: Amplitude-based kNN algorithm

efesafefsf

Quantum Computing & Qubits




Figure 1: Arbitrary two-dimensional qubit $|\psi\rangle$ visualized on the Bloch sphere¹

Most general form of a 2-D qubit:

$$|q\rangle = \alpha |0\rangle + \beta |1\rangle \quad (1)$$

where $\alpha, \beta \in \mathbb{C}$.

Can also be visualized in spherical polar coords on the unit or Bloch sphere as follows:

$$|q\rangle = \cos \frac{\theta}{2} |0\rangle + e^{i\phi} \sin \frac{\theta}{2} |1\rangle \quad (2)$$

where $0 \leq \theta \leq \pi$ and $0 \leq \phi \leq 2\pi$

¹Reprinted from Wikipedia, n.d., Retrieved September 7, 2016, from https://en.wikipedia.org/wiki/Bloch_Sphere. Copyright 2012 by Glosser.ca. Reprinted with permission.

- Approximately 2.5 quintillion (10^{18}) bytes of digital data are created every day¹
- Need for advanced algorithms that can make sense of data content, retrieve patterns and reveal correlations → Machine learning (ML)
- ML algorithms often involve
 - solving large systems of linear equations
 - inverting large matrices
 - distance computations
- Performing these computations on large data sets gets increasingly difficult²

¹IBM. (2016). What is big data? <https://www-01.ibm.com/software/data/bigdata/what-is-big-data.html>. (Accessed: 2016-09-08)

²Bekkerman, R., Bilenko, M., & Langford, J. (2011). Scaling up machine learning: Parallel and distributed approaches. Cambridge University Press.

1. ML involves manipulation of large vectors and matrices
 2. Quantum mechanics is about vectors \in complex Hilbert spaces
 3. Quantum computers are performing linear operations on qubits
- Hence, we can manipulate large vectors in parallel on quantum computers

So can we use QC to improve classical ML algorithms??

- Classical ML is a very practical topic
- BUT, QML has been of almost entirely theoretical nature

Quantum data encoding

There are two fundamentally different ways for state preparation:

Data encoded into qubits

k -dimensional probability vector requires $4k$ classical bits which are encoded one-to-one into $4k$ qubits, e.g.

$$\begin{pmatrix} 0.6 \\ 0.4 \end{pmatrix} * 10 \rightarrow \begin{pmatrix} 6 \\ 4 \end{pmatrix} \rightarrow \begin{pmatrix} 0110 \\ 0100 \end{pmatrix} \rightarrow n = 01100100 \rightarrow |n\rangle = |01100100\rangle$$

Data encoded into amplitudes

k -dimensional probability vector is encoded into $\log_2(k)$ qubits, e.g.

$$\begin{pmatrix} 0.6 \\ 0.4 \end{pmatrix} \rightarrow |n\rangle = \sqrt{0.6} |0\rangle + \sqrt{0.4} |1\rangle$$

Quantum data encoding

There are two fundamentally different ways for state preparation:

Data encoded into qubits

k -dimensional probability vector requires $4k$ classical bits which are encoded one-to-one into $4k$ qubits, e.g.

$$\begin{pmatrix} 0.6 \\ 0.4 \end{pmatrix} * 10 \rightarrow \begin{pmatrix} 6 \\ 4 \end{pmatrix} \rightarrow \begin{pmatrix} 0110 \\ 0100 \end{pmatrix} \rightarrow n = 01100100 \rightarrow |n\rangle = |01100100\rangle$$

Data encoded into amplitudes

k -dimensional probability vector is encoded into $\log_2(k)$ qubits, e.g.

$$\begin{pmatrix} 0.6 \\ 0.4 \end{pmatrix} \rightarrow |n\rangle = \sqrt{0.6} |0\rangle + \sqrt{0.4} |1\rangle$$

Classical k-nearest neighbour

- kNN is a non-parametric classifier
- k is a positive integer, usually chosen small

Given training data set: Given a new vector \tilde{x} (red star):

$$D_T = v_0, v_1, \dots, v_{10}$$

$$v_i \in \{A, B\}$$

- consider k nearest neighbours

- classify \tilde{x} , based on majority vote, as A or B




Figure 2: Visualization of a kNN classifier¹

¹Reprinted from GitHub, Burton de Wilde, Retrieved September 13, 2016, from <http://bdewilde.github.io/blog/blogger/2012/10/26/classification-of-hand-written-digits-3/>. Copyright 2012 by Burton de Wilde. Reprinted with permission.

The algorithm

$$\frac{1}{\sqrt{2M}} \sum_{m=1}^M (|0\rangle |\Psi_{\tilde{x}}(\star)\rangle + |1\rangle |\Psi_{x^m}\rangle) |y^m(\text{A or B})\rangle |m\rangle \quad (3)$$

where

$$|\Psi_{\tilde{x}}(\star)\rangle = \sum_{i=1}^N \tilde{x}_i |i\rangle \quad |\Psi_{x^m}\rangle = \sum_{i=1}^N x_i^m |i\rangle \quad (4)$$

$$\text{e.g.} \quad \begin{pmatrix} 0.6 \\ 0.4 \end{pmatrix} \rightarrow |n\rangle = \sqrt{0.6} |0\rangle + \sqrt{0.4} |1\rangle \quad (5)$$




Figure 3: Visualization of a kNN classifier¹

The algorithm

Applying the **Hadamard gate** interferes the input and the training vectors:

$$\frac{1}{2\sqrt{M}} \sum_{m=1}^M (|0\rangle[|\psi_{\tilde{x}}\rangle + |\psi_{x^m}\rangle] + |1\rangle[|\psi_{\tilde{x}}\rangle - |\psi_{x^m}\rangle]) |y^m(A \text{ or } B)\rangle |m\rangle \quad (6)$$

→ Perform **conditional measurement** on ancilla qubit.
Successful if $|0\rangle$ state is measured.

The algorithm

After successful conditional measurement, the state is proportional to

$$\frac{1}{2\sqrt{M}} \sum_{m=1}^M \sum_{i=1}^N (\tilde{x}_i + x_i^m) |0\rangle |i\rangle |y^m(\text{A or B})\rangle |m\rangle \quad (7)$$

Probability to measure class B:

$$p(|y^m\rangle = |1(\text{B})\rangle) = \sum_{m|y^m=1(\text{B})} 1 - \frac{1}{4M} |\tilde{x} - x^m|^2 \quad (8)$$

Overall algorithmic complexity

$O(\frac{1}{p_{acc}})$ where p_{acc} is the probability of measuring ancilla in the $|0\rangle$ state

Simple binary classification case




Figure 4: Simple binary classification problem of a quantum state

$$\frac{1}{\sqrt{2M}} \sum_{m=1}^M (|0\rangle |\Psi_{\tilde{x}}(\star)\rangle + |1\rangle |\Psi_{x^m}\rangle) |y^m(\text{A or B})\rangle |m\rangle \quad (9)$$

Procedure to load the input vector \tilde{x} :

$$|\Psi_0\rangle = \frac{1}{2} \sum_{m=1}^2 (|0\rangle |0\rangle + |1\rangle |0\rangle) |y^m\rangle |m\rangle \quad (10)$$

Apply controlled rotation ${}_0^1CR_y(\frac{\pi}{4})$ s.t.

$${}_0^1CR_y(\frac{\pi}{4}) |\Psi_0\rangle = |\Psi_1\rangle = \frac{1}{2} \sum_{m=1}^2 (|0\rangle |0\rangle + |1\rangle |\Psi_{\tilde{x}}\rangle) |y^m\rangle |m\rangle \quad (11)$$

Flip the ancilla qubit in the first register

$$(X \otimes \mathbb{1} \otimes \mathbb{1} \otimes \mathbb{1}) |\Psi_1\rangle = |\Psi_2\rangle = \frac{1}{2} \sum_{m=1}^2 (|0\rangle |\Psi_{\tilde{x}}\rangle + |1\rangle |0\rangle) |y^m\rangle |m\rangle \quad (12)$$

Implementation with IBM's quantum computer




Figure 5: IBM's quantum composer

- Accessible to the public
- Allows for ideal + real simulations
- 5 superconducting qubits
- 40 gates (39 gates + 1 measurement)

IBM's universal gate set



Figure 6: IBM's universal gate set

How can we implement the $\frac{1}{0}CR_y(\frac{\pi}{4})$ gate?

Controlled U gate




Figure 7: Controlled U-gate

Choose A,B,C and α s.t.




Figure 8: Decomposition of a controlled U-gate¹

$$e^{i\alpha} * A * X * B * X * C = U \quad \text{and} \quad A * B * C = \mathbb{1} \quad (13)$$

Need to solve the following equation¹

$$U = \begin{pmatrix} e^{i(\alpha - \frac{\beta}{2} - \frac{\delta}{2})} \cos \frac{\gamma}{2} & -e^{i(\alpha - \frac{\beta}{2} + \frac{\delta}{2})} \sin \frac{\gamma}{2} \\ e^{i(\alpha + \frac{\beta}{2} - \frac{\delta}{2})} \sin \frac{\gamma}{2} & e^{i(\alpha + \frac{\beta}{2} + \frac{\delta}{2})} \cos \frac{\gamma}{2} \end{pmatrix} \quad (14)$$

Overall algorithmic complexity

$O(\frac{1}{\rho_{acc}}) + O(k)$ where k is number of root finding iterations²

¹Nielsen, M. A., & Chuang, I. L. (2010). Quantum computation and quantum information. Cambridge University Press.

²Jat, R. N., & Ruhela, D. S. (2011). Comparative study of complexity of algorithms for iterative solution of non-linear equations. Journal of International Academy Of Physical Sciences, 15(4).

Problems with universal gate sets

In our case we need to find A, B, C and α for $\frac{1}{0}CR_Y(\frac{\pi}{4})$:

Using a root finding algorithm for non-linear equations we find:

$$\alpha = \pi; \quad \beta = 2\pi; \quad \delta = \frac{7}{8}\pi; \quad \gamma = 0 \quad (15)$$

Then,

$$A = R_z(\beta)R_y(\frac{\gamma}{2}) = R_z(2\pi) = \mathbb{1} \quad (16)$$

$$B = R_y(-\frac{\gamma}{2})R_z(-\frac{\delta + \beta}{2}) = R_z(-\frac{23}{16}\pi) = ??? \quad (17)$$

$$C = R_z(\frac{\delta - \beta}{2}) = R_z(-\frac{9}{16}\pi) = ??? \quad (18)$$

$$\begin{pmatrix} 1 & 0 \\ 0 & e^{i\alpha} \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ 0 & e^{i\pi} \end{pmatrix} = Z \quad (19)$$

The Solovay-Kitaev theorem

$$B = R_z\left(-\frac{23}{16}\pi\right) = ??? \quad (20)$$

$$C = R_z\left(-\frac{9}{16}\pi\right) = ??? \quad (21)$$

The Solovay-Kitaev theorem guarantees that given a set of single-qubit quantum gates which generates a dense subset of $SU(2)$, then that set is guaranteed to fill $SU(2)$ quickly.¹

→ **Hence, given any universal gate set it is possible to obtain good approximations to any desired gate.**


→ **But needs to be computed classically!**

¹Dawson, C. M., & Nielsen, M. A. (2005). The Solovay-Kitaev algorithm. arXiv preprint quant-ph/0505030.

The Solovay-Kitaev algorithm


Fowler distance¹:

$$\text{dist}(U, U_{\text{approx}}) = \sqrt{\frac{2 - |\text{tr}(U \cdot U_{\text{approx}}^\dagger)|}{2}} \quad (22)$$



¹Booth Jr, J. (2012). Quantum compiler optimizations. arXiv preprint arXiv:1206.3348.

The Solovay-Kitaev algorithm



$$d = 0.22739 \quad (23)$$

$$d = 0.15165 \quad (24)$$

$$d = 0.10722 \quad (25)$$

$$d = 0.02086 \quad (26)$$

$$d = 0.00156 \quad (27)$$

Figure 9: Various Fowler distances visualized on Bloch sphere

The Solovay-Kitaev algorithm

IBM's quantum computer needs **130ns for single-qubit gates** and **500ns for CNOT gates**.

IBM qubit decoherence times:

$49.5 \mu\text{s} \leq T_1 \leq 85.3 \mu\text{s}$ "amplitude damping"

$56.0 \mu\text{s} \leq T_2 \leq 139.7 \mu\text{s}$ "phase damping"

| Approx. Gate | Distance | Gate count | Execution time |
|--------------------------|----------|------------|----------------------------|
| $R_z(-\frac{23}{16}\pi)$ | 0.15165 | 25 | $\sim 3 \mu\text{s}$ |
| | 0.10722 | 109 | $\sim 14 \mu\text{s}$ |
| | 0.02086 | 2997 | $\sim 390 \mu\text{s}$ |
| | 0.01494 | 14721 | $\sim 1914 \mu\text{s}$ |
| | 0.003327 | 74009 | $\sim 9621 \mu\text{s}$ |
| | 0.001578 | 370813 | $\sim 48\,206 \mu\text{s}$ |

Table 1: SK algorithm results

Adding complexities

Executing the SK algorithm adds to our overall algorithmic complexity:

Overall algorithmic complexity

$O(\frac{1}{p_{acc}}) + O(k) + O(m * \log^{2.71}(\frac{m}{\epsilon}))$ for ϵ -approximations of m gates¹

Due to state preparation we went from

$$O(\frac{1}{p_{acc}}) \tag{28}$$

suddenly to

$$O(m * \log^{2.71}(\frac{m}{\epsilon})) \tag{29}$$

where m is the number of gates that need approximation to ϵ -accuracy

¹Dawson, C. M., & Nielsen, M. A. (2005). The Solovay-Kitaev algorithm. arXiv preprint quant-ph/0505030.

Currently impossible to implement the quantum algorithm on IBM's quantum computer! \rightarrow can only simulate it with i.e. Liqui $\lvert\rangle$

In Liqui $\lvert\rangle$ we can directly implement the controlled R_y rotation!

Conclusion

Summary

- Initial state preparation is non-trivial! (even for very simple examples)
- In this case, state preparation dominates the overall algorithmic complexity
- Solovay-Kitaev yields long gate sequences for good approximations
- Some universal gate sets are only useful when combined with long qubit lifetimes
- **Need for better quantum compiling and more general state preparation algorithms!**

- Complexity analysis of the qubit-based kNN algorithm
- Classification of gaussian probability distributions
- Implementing more general state preparation algorithms
- Waiting for IBM QASM 2.0 ...

References

Bekkerman, R., Bilenko, M., & Langford, J. (2011). Scaling up machine learning: Parallel and distributed approaches. Cambridge University Press.

Booth Jr, J. (2012). Quantum compiler optimizations. arXiv preprint arXiv:1206.3348.

Dawson, C. M., & Nielsen, M. A. (2005). The Solovay-Kitaev algorithm. arXiv preprint quant-ph/0505030.

IBM. (2016). What is big data?

<https://www-01.ibm.com/software/data/bigdata/what-is-big-data.html>. (Accessed: 2016-09-08)

Jat, R. N., & Ruhela, D. S. (2011). Comparative study of complexity of algorithms for iterative solution of non-linear equations. Journal of International Academy Of Physical Sciences, 15(4).

Nielsen, M. A., & Chuang, I. L. (2010). Quantum computation and quantum information. Cambridge University Press.

Questions?

Backup slide: Qubit decoherence times

We expect exponential decay:

$$e^{t/T_i} \quad (30)$$

T_1 : Longitudinal coherence time (amplitude damping)

- Prepare $|0\rangle$ state
- Apply the X (NOT) gate s.t. qubit is in $|1\rangle$ state
- Wait for time t
- Measure the probability of being in $|1\rangle$ state

T_2 : Transversal coherence time (phase damping)

- Prepare $|0\rangle$ state
- Apply Hadamard $\rightarrow \frac{|0\rangle + |1\rangle}{\sqrt{2}}$
- Wait for time t
- Apply Hadamard again
- Measure the probability of being in $|0\rangle$ state

We expect this probability to go to 0.5 \rightarrow qubit lost quantum behaviour

Backup Slide II: Experimental realizations

Only few experimental verifications of QML algorithms:

- Li, Liu, Xu, and Du (2015) successfully distinguished a handwritten six from a nine using a quantum support vector machine on a four-qubit nuclear magnetic resonance test bench¹
- Cai et al. (2015) were first to experimentally demonstrate quantum machine learning on a photonic QC and showed that the distance between two vectors and their inner product can indeed be computed quantum mechanically²
- Rist et al. (2015) solved a learning parity problem with five superconducting qubits and found that a quantum advantage can already be observed in non error-corrected systems³

¹Li, Z., Liu, X., Xu, N., & Du, J. (2015). Experimental realization of a quantum support vector machine. Physical Review Letters, 114 (14), 15. doi: 10.1103/PhysRevLett.114.140504

²Cai, X. D., Wu, D., Su, Z. E., Chen, M. C., Wang, X. L., Li, L., . . . Pan, J. W. (2015). Entanglement- based machine learning on a quantum computer.

Machine Learning

Machine learning can be subdivided into three major fields.

Supervised ML

- Based on *input* and *output* data

"I know how to classify this data but I need the algorithm to do the computations for me."

Unsupervised ML

- Based on *input* data only

"I have no clue how to classify this data, can the algorithm create a classifier for me?"

Reinforcement learning

- Based on *input* data only

"I have no clue how to classify this data, can the algorithm classify this data and I'll give it a reward if it's correct or I'll punish it if it's not."

Machine Learning

Machine learning can be subdivided into three major fields.

Supervised ML

- Based on *input* and *output* data

"I know how to classify this data but I need the algorithm to do the computations for me."

Unsupervised ML

- Based on *input* data only

"I have no clue how to classify this data, can the algorithm create a classifier for me?"

Reinforcement learning

- Based on *input* data only

"I have no clue how to classify this data, can the algorithm classify this data and I'll give it a reward if it's correct or I'll punish it if it's not."

Liqui|⟩ simulations: Taking it further





Figure 10: Representation of hamming distance on 3D cube

Liqui|⟩ simulations: Taking it further

Applying the following matrix

$$\begin{pmatrix} \sqrt{\delta} & 1 - \sqrt{\delta} \\ 1 - \sqrt{\delta} & -\sqrt{\delta} \end{pmatrix} \quad (31)$$

to all qubits in the data register leads to a gaussian distribution over the "Hamming distance" cube:



Liqui|> simulations: Taking it further




Figure 12: Representation of gaussian diffusion on 3D cube

Liqui|⟩ simulations: Taking it further




Figure 13: Representation of gaussian diffusion on 3D cube



IBM.

What is big data?

<https://www-01.ibm.com/software/data/bigdata/what-is-big-data.html>, 2016.

Accessed: 2016-09-08.

Qubit-based kNN quantum algorithm

Typography

The theme provides sensible defaults to
`\emph{emphasize}` text, `\alert{accent}` parts
or show `\textbf{bold}` results.

becomes

The theme provides sensible defaults to *emphasize* text, **accent** parts or
show **bold** results.

Font feature test

- Regular
- *Italic*
- SMALLCAPS
- **Bold**
- **Bold Italic**
- **Bold SmallCaps**
- Monospace
- *Monospace Italic*
- Monospace Bold
- *Monospace Bold Italic*

Lists

Items

- Milk
- Eggs
- Potatos

Enumerations

1. First,
2. Second and
3. Last.

Descriptions

PowerPoint Meeh.
Beamer Yeeeha.

- This is important

Animation

- This is important
- Now this

Animation

- This is important
- Now this
- And now this

Animation

- This is really important
- Now this
- And now this

Figures




Figure 14: Rotated square from texample.net.

Table 2: Largest cities in the world (source: Wikipedia)

| City | Population |
|-------------|------------|
| Mexico City | 20,116,842 |
| Shanghai | 19,210,000 |
| Peking | 15,796,450 |
| Istanbul | 14,160,467 |

Blocks

Three different block environments are pre-defined and may be styled with an optional background color.

Default

Block content.

Alert

Block content.

Example

Block content.

Default

Block content.

Alert


Block content.

Example


Block content.

$$e = \lim_{n \rightarrow \infty} \left(1 + \frac{1}{n}\right)^n$$

Line plots



Bar charts



Veni, Vidi, Vici